

# Dueling Bandit Review

Bowen Xu

Sep 2022

# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner
- 5 Dependent Arms
- 6 Conclusion

# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner
- 5 Dependent Arms
- 6 Conclusion

Drawback of Conventional MAB problem:

- When payoff is a **relative comparison** result rather than an **absolute value**, it is difficult to apply conventional MAB Algorithm.

What if Rewards aren't Directly Measurable?

Some scenarios where the traditional MAB algorithm cannot be applied:

- User-perceived quality of a set of retrieval results
- taste of food
- product attractiveness

Characteristics of dueling bandit algorithm:

- Only **binary feedback** about the **relative reward** of two chosen strategies is available
- Pairwise comparison is made in this problem.
- Dueling Bandit Problem applies where a system must adapt interactively to specific user bases

Suitable application scenarios:

- Search-engine user prefers ranking  $r_1$  over  $r_2$  for a given query
- Online Advertising
- Recommender Systems: (e.g. Restaurant recommend app)
- Ongoing work: Personalized Clinical Treatment.



# Application of Dueling Bandit Algorithm




(a) search engine



(b) online advertisement




(c) recommend system



(d) personalized clinical treatment

# Recommend System

- Experiment 1: Pepsi or Coca
- Model: A 2-arm dueling bandit problem



Compare the Coca and Pepsi cola and we will find that the Pepsi cola is more preferred than Coca cola relatively.

# Search Engine

- Experiment 2: Rank the customers' preference of websites
- Model: A Dueling bandit problem

## [Web-Page Summarization Using Clickthrough Data - Microsoft Research](#)

By Jian-Tao Sun, Dou Shen, HuaJun Zeng, Qiang Yang, Yuchang Lu and Zheng Chen. In: Proceedings of the 28th Annual International ACM SIGIR Conference, August 2005. The ...  
[research.microsoft.com/apps/pubs/default.aspx?id=69202](#) · Mark as spam

## [Optimizing Search Engines using Clickthrough Data](#)

Optimizing Search Engines using Clickthrough Data Thorsten Joachims Cornell University  
Computer Science Ithaca, NY 14853 USA [tj@cs.cornell.edu](mailto:tj@cs.cornell.edu) ABSTRACT ...  
[www.cs.cornell.edu/People/tj/publications/joachims\\_02c.pdf](#) · PDF file · Mark as spam



## [Clickthrough Data](#)

This page shows one keyword best matching your query, you can find other results here.  
[academic.research.microsoft.com/Search.aspx?query=Clickthrough+data](#) · Mark as spam

## [Smoothing clickthrough data for web search ranking](#)

Incorporating features extracted from clickthrough data (called clickthrough features) has been demonstrated to significantly improve the performance of ranking models for ...  
[academic.research.microsoft.com/Paper/5432909.aspx](#) · Mark as spam

## [CiteSeerX — Smoothing Clickthrough Data for Web Search Ranking](#)

CiteSeerX - Document Details (Isaac Council, Lee Giles): Incorporating features extracted from clickthrough data (called clickthrough features) has been demonstrated to ...  
[citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.150.2058](#) · Mark as spam

## [CiteSeerX — How Does Clickthrough Data Reflect Retrieval Quality?](#)

@MISC(Radlinski\_howdoes, author = {Filip Radlinski and Madhu Kurup and Thorsten Joachims}, title = {How Does Clickthrough Data Reflect Retrieval Quality?}, year = {}  
[citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.147.454](#) · Mark as spam

# Search Engine: Interleave

Interleave the two reference ranking into the presented ranking.

Method: At each time, a coin flip decides which captain can choose his next teammate.




# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner
- 5 Dependent Arms
- 6 Conclusion

# Dueling Mechanism

Each pair duels until statistical significance



# Dueling Mechanism



Interleave A vs B



...

	Left wins	Right wins
A vs B	0	1
A vs C	0	0
B vs C	0	0

[From Yisong Yue]

# Dueling Mechanism

Ranking A	Ranking B
<ol style="list-style-type: none"> <li>1. Nepal Valley - The valley for writing.</li> <li>2. Nepal Valley - The valley for writing.</li> <li>3. Nepal Valley - The valley for writing.</li> <li>4. Nepal Valley - The valley for writing.</li> <li>5. Nepal Valley - The valley for writing.</li> <li>6. Nepal Valley - The valley for writing.</li> <li>7. Nepal Valley - The valley for writing.</li> <li>8. Nepal Valley - The valley for writing.</li> <li>9. Nepal Valley - The valley for writing.</li> <li>10. Nepal Valley - The valley for writing.</li> </ol>	<ol style="list-style-type: none"> <li>1. Nepal Valley - The valley for writing.</li> <li>2. Nepal Valley - The valley for writing.</li> <li>3. Nepal Valley - The valley for writing.</li> <li>4. Nepal Valley - The valley for writing.</li> <li>5. Nepal Valley - The valley for writing.</li> <li>6. Nepal Valley - The valley for writing.</li> <li>7. Nepal Valley - The valley for writing.</li> <li>8. Nepal Valley - The valley for writing.</li> <li>9. Nepal Valley - The valley for writing.</li> <li>10. Nepal Valley - The valley for writing.</li> </ol>

Interleave A vs C



...

	Left wins	Right wins
A vs B	0	1
A vs C	0	<b>1</b>
B vs C	0	0

[From Yisong Yue]



# Dueling Mechanism



### Interleave B vs C



	Left wins	Right wins
A vs B	0	1
A vs C	0	1
B vs C	0	1

[From Yisong Yue]

# Dueling Mechanism



Interleave A vs C



...

	Left wins	Right wins
A vs B	0	1
A vs C	1	1
B vs C	0	1

[From Yisong Yue]



## Dueling Bandits Problem



**Goal:** Maximize total user utility

**Exploit:** run C  
(interleave C with itself)

**Explore:** interleave A vs B

**Best:** A  
(interleave A with itself)

How to interact optimally?

	Left wins	Right wins
A vs B	0	1
A vs C	1	1
B vs C	0	1

[From Yisong Yue]

# Problem Setting

## Dueling Bandits:

- Process: At each time step:  $t = 1, 2, \dots, T$ 
  - The algorithm chooses a pair of actions  $a_i, a_j$  from  $K$  available actions.
  - The world provides (independent stochastic) preference feedback of which action is more preferred.
  - The preference feedback satisfy the probability of  $P(a_i > a_j) = P_{ij}$

# Problem Setting

## Dueling Bandits:

- Process: At each time step:  $t = 1, 2, \dots, T$ 
  - The algorithm chooses a pair of actions  $a_i, a_j$  from  $K$  available actions.
  - The world provides (independent stochastic) preference feedback of which action is more preferred.
  - The preference feedback satisfy the probability of  $P(a_i > a_j) = P_{ij}$
- Regret:  $R(T) = \sum_{t=1}^T r(t)$ 
  - $r = \Delta_{1i} + \Delta_{1j}$ ,  $\Delta_{ij} = P_{ij} - 0.5$

# Problem Setting

## Dueling Bandits:

- Process: At each time step:  $t = 1, 2, \dots, T$ 
  - The algorithm chooses a pair of actions  $a_i, a_j$  from  $K$  available actions.
  - The world provides (independent stochastic) preference feedback of which action is more preferred.
  - The preference feedback satisfy the probability of  $P(a_i > a_j) = P_{ij}$
- Regret:  $R(T) = \sum_{t=1}^T r(t)$ 
  - $r = \Delta_{1i} + \Delta_{1j}$ ,  $\Delta_{ij} = P_{ij} - 0.5$
- Goal: Minimize the cumulative regret:  $R(T)$  from pulling the suboptimal arms.

# Problem Setting: Winner Setting

- Condorcet Winner: the arm beat all other arms with probability of 0.5 ( $P_{ij} > 0.5$ )
  - e.g. IF, BTM, Sparring, RUCB
- Copeland Winner: the arm beat the most other arms with probability of 0.5.
  - e.g. CCB, RCB, RMED, DTS

## Two Styles of Algorithm Design:

- Asymmetric Algorithms: choosing a reference arm and a exploration arm.
  - e.g. IF, BtM, SAVAGE, Doubler, RUCB, MergeRUCB, RCS, and DTS.
- Symmetric Algorithms: treats the choice of the two arms symmetrically.
  - e.g. Sparring, Self-Sparring.



# Problem Setting

Basic Assumption:

- Ordered Arms: We can relabel arms as  $P_{ij} > 0.5$  for all  $i, j$ .

# Problem Setting

Basic Assumption:

- Ordered Arms: We can relabel arms as  $P_{ij} > 0.5$  for all  $i, j$ .
- Stochastic Triangle Inequality (STI):  $\Delta_{ik} \leq \Delta_{ij} + \Delta_{jk} (i \leq j \leq k)$ .


# Problem Setting

## Basic Assumption:


- Ordered Arms: We can relabel arms as  $P_{ij} > 0.5$  for all  $i, j$ .
- Stochastic Triangle Inequality (STI):  $\Delta_{ik} \leq \Delta_{ij} + \Delta_{jk} (i \leq j \leq k)$ .
- Transitivity Conditions:
  - Strong Stochastic Transitivity (SST):  $\Delta_{ik} \geq \max\{\Delta_{ij}, \Delta_{jk}\}$
  - Relaxed Stochastic Transitivity (RST):  $\gamma \Delta_{ik} \geq \max\{\Delta_{ij}, \Delta_{jk}\}$

## Strong Stochastic Transitivity

$$\Delta_{ik} \geq \max\{\Delta_{ij}, \Delta_{jk}\}.$$

Monotonic 

	A	B	C	D	E	F
A	0	0.03	0.04	0.06	0.10	0.11
B	-0.03	0	0.03	0.05	0.08	0.11
C	-0.04	-0.03	0	0.04	0.07	0.09
D	-0.06	-0.05	-0.04	0	0.05	0.07
E	-0.10	-0.08	-0.07	-0.05	0	0.03
F	-0.11	-0.11	-0.09	-0.07	-0.03	0

 Monotonic

## Strong Triangle Inequality

$$\Delta_{ik} \leq \Delta_{ij} + \Delta_{jk}.$$

**Red**  $\leq$  **Blue** + **Green**

	A	B	C	D	E	F
A	0	0.03	0.04	0.06	0.10	0.11
B	-0.03	0	0.03	0.05	0.08	0.11
C	-0.04	-0.03	0	0.04	0.07	0.09
D	-0.06	-0.05	-0.04	0	0.05	0.07
E	-0.10	-0.08	-0.07	-0.05	0	0.03
F	-0.11	-0.11	-0.09	-0.07	-0.03	0

# Development (Main Algorithm)

- Intereaved Filter [Yue et al., 2009]
- Beat the Mean [Yue, Joachims., 2011]
- SAVAGE [Urvoy et al., 2013]
- Sparring [Alion et al., 2014]
- RMED [Komiyama et al., 2015]
- RUCB [Zoghi et al., 2014; 2015]
- CCB, SCB [Masrour Zoghi et al., 2015]
- DTS [Wu, Liu, 2016]
- SelfSparring [Yanan Sui et al., 2017]
- .....

Two main Model Assumption:

- Condorcet Winner
- Copeland Winner [Zoghi et al., 2015]

Other Model Assumption:

- Borda Winner [Jamieson et al., 2015]
- Von Neuman Winner [Dudik et al., 2015]
- General Tournament Solutions [Ramamohan et al. 2016]

# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner**
- 4 Copeland Winner
- 5 Dependent Arms
- 6 Conclusion



# Interleaved Filter (IF)

The basic algorithm of K-arm Dueling Bandits problem.

- Motivation: solve the problem where absolute rewards have no natural scale or are difficult to measure.  
(K-arm Bandit problem)

# Interleaved Filter (IF)

The basic algorithm of K-arm Dueling Bandits problem.

- Motivation: solve the problem where absolute rewards have no natural scale or are difficult to measure.  
(K-arm Bandit problem)
- Algorithm: IF1, IF2

# IF Algorithm: IF1

- Idea: combine estimate  $\hat{P}$  with confidence interval  $\hat{C}$ .  
 $\hat{P}_{b',b} = P(b' > b), \hat{C}_t = (\hat{P}_t - c_t, \hat{P}_t + c_t), c_t = \sqrt{\log(1/\delta)/t}$

# IF Algorithm: IF1

- Idea: combine estimate  $\hat{P}$  with confidence interval  $\hat{C}$ .  
 $\hat{P}_{b',b} = P(b' > b), \hat{C}_t = (\hat{P}_t - c_t, \hat{P}_t + c_t), c_t = \sqrt{\log(1/\delta)/t}$
- Assumption: SST, finite horizon

# IF Algorithm: IF1

- Idea: combine estimate  $\hat{P}$  with confidence interval  $\hat{C}$ .  
 $\hat{P}_{b',b} = P(b' > b)$ ,  $\hat{C}_t = (\hat{P}_t - c_t, \hat{P}_t + c_t)$ ,  $c_t = \sqrt{\log(1/\delta)/t}$
- Assumption: SST, finite horizon
- Regret:  $E(R_T^{IF1}) = O(\frac{K \log K}{\Delta_{1,2}} \log T)$  ( $\Delta_{1,2}$  means the distinguishability between the two best bandits)

# Algorithm: IF1

---

**Algorithm 2** Interleaved Filter 1 (IF1)

---

```
1: Input:  $T, \mathcal{B} = \{b_1, \dots, b_K\}$ 
2:  $\delta \leftarrow 1/(TK^2)$ 
3: Choose  $\hat{b} \in \mathcal{B}$  randomly
4:  $W \leftarrow \{b_1, \dots, b_K\} \setminus \{\hat{b}\}$ 
5:  $\forall b \in W$ , maintain estimate  $\hat{P}_{\hat{b},b}$  of  $P(\hat{b} > b)$ 
6:  $\forall b \in W$ , maintain  $1 - \delta$  confidence interval  $\hat{C}_{\hat{b},b}$  of  $\hat{P}_{\hat{b},b}$ 
7: while  $W \neq \emptyset$  do
8:   for  $b \in W$  do
9:     compare  $\hat{b}$  and  $b$ 
10:    update  $\hat{P}_{\hat{b},b}, \hat{C}_{\hat{b},b}$ 
11:   end for
12:   while  $\exists b \in W$  s.t.  $(\hat{P}_{\hat{b},b} > 1/2 \wedge 1/2 \notin \hat{C}_{\hat{b},b})$  do
13:      $W \leftarrow W \setminus \{b\}$ 
14:   end while
15:   if  $\exists b' \in W$  s.t.  $(\hat{P}_{\hat{b},b'} < 1/2 \wedge 1/2 \notin \hat{C}_{\hat{b},b'})$  then
16:      $\hat{b} \leftarrow b', W \leftarrow W \setminus \{b'\}$  //new round
17:      $\forall b \in W$ , reset  $\hat{P}_{\hat{b},b}$  and  $\hat{C}_{\hat{b},b}$ 
18:   end if
19: end while
20:  $\hat{T} \leftarrow$  Total Comparisons Made
21: return  $(\hat{b}, \hat{T})$ 
```

---

# IF Algorithm: IF2

- Idea: IF1 + **pruning**.

# IF Algorithm: IF2

- Idea: IF1 + **pruning**.
- Assumption: SST, finite horizon



# IF Algorithm: IF2

- Idea: IF1 + **pruning**.
- Assumption: SST, finite horizon
- Regret:  $E(R_T^{IF2}) = O(\frac{K}{\Delta_{1,2}} \log T)$

# Algorithm: IF2

---

**Algorithm 3** Interleaved Filter 2 (IF2)

---

```
1: Input:  $T, \mathcal{B} = \{b_1, \dots, b_K\}$ 
2:  $\delta \leftarrow 1/(TK^2)$ 
3: Choose  $\hat{b} \in \mathcal{B}$  randomly
4:  $W \leftarrow \{b_1, \dots, b_K\} \setminus \{\hat{b}\}$ 
5:  $\forall b \in W$ , maintain estimate  $\hat{P}_{\hat{b},b}$  of  $P(\hat{b} > b)$ 
6:  $\forall b \in W$ , maintain  $1 - \delta$  confidence interval  $\hat{C}_{\hat{b},b}$  of  $\hat{P}_{\hat{b},b}$ 
7: while  $W \neq \emptyset$  do
8:   for  $b \in W$  do
9:     compare  $\hat{b}$  and  $b$ 
10:    update  $\hat{P}_{\hat{b},b}, \hat{C}_{\hat{b},b}$ 
11:   end for
12:   while  $\exists b \in W$  s.t.  $(\hat{P}_{\hat{b},b} > 1/2 \wedge 1/2 \notin \hat{C}_{\hat{b},b})$  do
13:      $W \leftarrow W \setminus \{b\}$ 
14:   end while
15:   if  $\exists b' \in W$  s.t.  $(\hat{P}_{\hat{b},b'} < 1/2 \wedge 1/2 \notin \hat{C}_{\hat{b},b'})$  then
16:     while  $\exists b \in W$  s.t.  $\hat{P}_{\hat{b},b} > 1/2$  do
17:        $W \leftarrow W \setminus \{b\}$  //pruning
18:     end while
19:      $\hat{b} \leftarrow b', W \leftarrow W \setminus \{b'\}$  //new round
20:      $\forall b \in W$ , reset  $\hat{P}_{\hat{b},b}$  and  $\hat{C}_{\hat{b},b}$ 
21:   end if
22: end while
23:  $\hat{T} \leftarrow$  Total Comparisons Made
24: return  $(\hat{b}, \hat{T})$ 
```

# Interleaved Filter

Reference: The K-armed Dueling bandits problem, 2009

# Beat the Mean (BTM)

- motivation: Extend the Dueling Bandits Problem to a **relaxed setting (RST)** where preference magnitudes can violate transitivity

# Beat the Mean (BTM)

- motivation: Extend the Dueling Bandits Problem to a **relaxed setting (RST)** where preference magnitudes can violate transitivity
- Idea: Compare the fewest recorded comparison bandit with the mean bandit (random sample to obtain the mean)

# Beat the Mean (BTM)

- motivation: Extend the Dueling Bandits Problem to a **relaxed setting (RST)** where preference magnitudes can violate transitivity
- Idea: Compare the fewest recorded comparison bandit with the mean bandit (random sample to obtain the mean)
- Assumption: **RST**, ordered arms

# Beat the Mean (BTM)

- motivation: Extend the Dueling Bandits Problem to a **relaxed setting (RST)** where preference magnitudes can violate transitivity
- Idea: Compare the fewest recorded comparison bandit with the mean bandit (random sample to obtain the mean)
- Assumption: **RST**, ordered arms
- Setting: PAC(confident near-optimal bandit), online

# Beat the Mean (BTM)

---

**Algorithm 1** BEAT-THE-MEAN

---

```
1: Input:  $\mathcal{B} = \{b_1, \dots, b_K\}$ ,  $N$ ,  $T$ ,  $c_{\delta, \gamma}(\cdot)$ 
2:  $W_1 \leftarrow \{b_1, \dots, b_K\}$  //working set of active bandits
3:  $\ell \leftarrow 1$  //num rounds
4:  $\forall b \in W_\ell, n_b \leftarrow 0$  //num comparisons
5:  $\forall b \in W_\ell, w_b \leftarrow 0$  //num wins
6:  $\forall b \in W_\ell, \hat{P}_b \equiv w_b/n_b$ , or  $1/2$  if  $n_b = 0$ 
7:  $n^* \equiv \min_{b \in W_\ell} n_b$ 
8:  $c^* \equiv c_{\delta, \gamma}(n^*)$ , or  $1$  if  $n^* = 0$  //confidence radius
9:  $t \leftarrow 0$  //total number of iterations
10: while  $|W_\ell| > 1$  and  $t < T$  and  $n^* < N$  do
11:    $b \leftarrow \operatorname{argmin}_{b \in W_\ell} n_b$  //break ties randomly
12:   select  $b' \in W_\ell$  at random, compare  $b$  vs  $b'$ 
13:   if  $b$  wins,  $w_b \leftarrow w_b + 1$ 
14:    $n_b \leftarrow n_b + 1$ 
15:    $t \leftarrow t + 1$ 
16:   if  $\min_{b' \in W_\ell} \hat{P}_{b'} + c^* \leq \max_{b \in W_\ell} \hat{P}_b - c^*$  then
17:      $b' \leftarrow \operatorname{argmin}_{b \in W_\ell} \hat{P}_b$ 
18:      $\forall b \in W_\ell$ , delete comparisons with  $b'$  from  $w_b, n_b$ 
19:      $W_{\ell+1} \leftarrow W_\ell \setminus \{b'\}$  //update working set
20:      $\ell \leftarrow \ell + 1$  //new round
21:   end if
22: end while
23: return  $\operatorname{argmax}_{b \in W_\ell} \hat{P}_b$ 
```

---



# Beat the Mean (BTM)

Regret:

- $(\epsilon - \delta)$ -PAC:  $O(KN) = O(\frac{K\gamma^6}{\epsilon^2} \log \frac{KN}{\delta})$
- Online:  $O(\sum_{l=1}^{K-1} \min\{\frac{\gamma^7}{\epsilon_l}, \frac{\gamma^5 \epsilon_l}{\epsilon_*^2}\} \log T) = O(\frac{\gamma^7 K}{\epsilon_*} \log T)$

# Another Perspective

Borda Score:  $\sum_j \frac{1}{K} p_{ij}$

Theorem:

- The Borda score of the Condorcet winner is always greater than or equal to 0.5.
- The Condorcet winner remains optimal after removing other bandits.

Idea: Keep eliminating Borda losers, eventually the Condorcet winner.

# Beat the Mean

Reference: Beat the Mean Bandit, 2011

Utility-based dueling bandits problem

- Average utility:  $U_t^{av} = (u_t + v_t)/2$

Utility-based dueling bandits problem

- Average utility:  $U_t^{av} = (u_t + v_t)/2$
- Reward(utility) not observed

## Utility-based dueling bandits problem

- Average utility:  $U_t^{av} = (u_t + v_t)/2$
- Reward(utility) not observed
- Link function:  $\Pr[b_t = 1 | (u_t, v_t)] = \Phi(u_t, v_t)$

- Motivation: Reducing the Dueling Bandits problem to the conventional (stochastic) **Multi-Armed Bandits** problem

- Motivation: Reducing the Dueling Bandits problem to the conventional (stochastic) **Multi-Armed Bandits** problem
- Idea: View the dueling bandit as the dueling of two arms with different strategy (left and right)



- Motivation: Reducing the Dueling Bandits problem to the conventional (stochastic) **Multi-Armed Bandits** problem
- Idea: View the dueling bandit as the dueling of two arms with different strategy (left and right)
- Algorithm: Doubler, MultiSBM, Sparring

# UBDB Algorithm: Doubler

Structured: Large or possibly infinite set of arms  $X$

- left arm: random choice
- right arm: SBM

Regret (SBM is UCB):  $O(H \log^2 T)$ , where  $H = \sum_{i=2}^K \Delta_i^{-1}$

# UBDB Algorithm: Doubler

---

**Algorithm 2 (Doubler):** Reduction for finite and infinite  $X$  with internal structure.

---

```
1:  $S \leftarrow$  new SBM over  $X$ 
2:  $\mathcal{L} \leftarrow$  an arbitrary singleton in  $X$ 
3:  $i \leftarrow 1, t \leftarrow 1$ 
4: while true do
5:   reset( $S$ )
6:   for  $j = 1 \dots 2^i$  do
7:     choose  $x_t$  uniformly from  $\mathcal{L}$ 
8:      $y_t \leftarrow$  advance( $S$ )
9:     play  $(x_t, y_t)$ , observe choice  $b_t$ 
10:    feedback( $S, b_t$ )
11:     $t \leftarrow t + 1$ 
12:   end for
13:    $\mathcal{L} \leftarrow$  the multi-set of arms played as  $y_t$  in the last
   for-loop
14:    $i \leftarrow i + 1$ 
15: end while
```

---

# UBDB Algorithm: MultiSBM

Unstructured: The elements of  $X$  typically have no structure

- left arm: arm of the last round
- right arm: SBM

Regret (SBM is UCB):  $O(H\alpha \ln T + H\alpha(K \ln K + K \ln \ln T - \sum_{x \neq x^*} \ln \Delta_x))$

# UBDB Algorithm: MultiSBM

---

**Algorithm 3 (MultiSBM):** Reduction for unstructured finite  $X$  by using  $K$  SBMs in parallel.

---

```
1: For all  $x \in X$ :  $S_x \leftarrow$  new SBM over  $X$ , reset( $S_x$ )
2:  $y_0 \leftarrow$  arbitrary element of  $X$ 
3:  $t \leftarrow 1$ 
4: while true do
5:    $x_t \leftarrow y_{t-1}$ 
6:    $y_t \leftarrow$  advance( $S_{x_t}$ )
7:   play ( $x_t, y_t$ ), observe choice  $b_t$ 
8:   feedback( $S_{x_t}, b_t$ )
9:    $t \leftarrow t + 1$ 
10: end while
```

---

# UBDB Algorithm: Sparring

- left arm: SBM1
- right arm: SBM2

Upper bound of Regret: a constant times the regret of the two SBMs

---

**Algorithm 4 (Sparring):** Reduction to two SBMs.

---

```
1:  $S_L, S_R \leftarrow$  two new SBMs over  $X$ 
2:  $\text{reset}(S_L), \text{reset}(S_R), t \leftarrow 1$ 
3: while true do
4:    $x_t \leftarrow \text{advance}(S_L); y_t \leftarrow \text{advance}(S_R)$ 
5:   play  $(x_t, y_t)$ , observe choice  $b_t \in \{0, 1\}$ 
6:   feedback( $S_L, \mathbf{1}_{b_t=0}$ ); feedback( $S_R, \mathbf{1}_{b_t=1}$ )
7:    $t \leftarrow t + 1$ 
8: end while
```

---

# Relative UCB (RUCB)

- Motivation: Extends the **Upper Confidence Bound (UCB)** algorithm in relative setting to apply dueling bandit



# Relative UCB (RUCB)

- Motivation: Extends the **Upper Confidence Bound (UCB)** algorithm in relative setting to apply dueling bandit
- Idea: Implement two comparisons:
  - optimistically:  $u_{uj}$  (line 6)
  - pessimistically:  $u_{jc}$  (line 13)

objective: avoid auto-comparisons (self-comparisons obtains no information) until is great certainty of Condorcet winner.

# Relative UCB (RUCB)

- Motivation: Extends the **Upper Confidence Bound (UCB)** algorithm in relative setting to apply dueling bandit
- Idea: Implement two comparisons:
  - optimistically:  $u_{uj}$  (line 6)
  - pessimistically:  $u_{jc}$  (line 13)

objective: avoid auto-comparisons (self-comparisons obtains no information) until is great certainty of Condorcet winner.

# Relative UCB (RUCB)

- Motivation: Extends the **Upper Confidence Bound (UCB)** algorithm in relative setting to apply dueling bandit
  - Idea: Implement two comparisons:
    - optimistically:  $u_{uj}$  (line 6)
    - pessimistically:  $u_{jc}$  (line 13)
- objective: avoid auto-comparisons (self-comparisons obtains no information) until is great certainty of Condorcet winner.
- Regret:  $O(K^2 + K \log T)$

# Relative UCB (RUCB)

- Strengths: RUCB achieves high-probability regret bounds while making minimal assumptions other than assuming a Condorcet winner.
- Weaknesses: Compare Condorcet winner with itself both optimistically and pessimistically, where it is overly prudent.

# RUCB Algorithm

---

**Algorithm 1** Relative Upper Confidence Bound

---

**Input:**  $\alpha > \frac{1}{2}$ ,  $T \in \{1, 2, \dots\} \cup \{\infty\}$

- 1:  $\mathbf{W} = [w_{ij}] \leftarrow \mathbf{0}_{K \times K}$  // 2D array of wins:  $w_{ij}$  is the number of times  $a_i$  beat  $a_j$
- 2:  $\mathcal{B} = \emptyset$
- 3: **for**  $t = 1, \dots, T$  **do**
- 4:  $\mathbf{U} := [u_{ij}] = \frac{\mathbf{W}}{\mathbf{W} + \mathbf{W}^T} + \sqrt{\frac{\alpha \ln t}{\mathbf{W} + \mathbf{W}^T}}$  // All operations are element-wise;  $\frac{x}{0} := 1$  for any  $x$ .
- 5:  $u_{ii} \leftarrow \frac{1}{2}$  for each  $i = 1, \dots, K$ .
- 6:  $\mathcal{C} \leftarrow \{a_c \mid \forall j : u_{cj} \geq \frac{1}{2}\}$ .
- 7: If  $\mathcal{C} = \emptyset$ , then pick  $c$  randomly from  $\{1, \dots, K\}$ .
- 8:  $\mathcal{B} \leftarrow \mathcal{B} \cup \mathcal{C}$ .
- 9: If  $|\mathcal{C}| = 1$ , then  $\mathcal{B} \leftarrow \mathcal{C}$  and  $a_c$  to be the unique element in  $\mathcal{C}$ .
- 10: **if**  $|\mathcal{C}| > 1$  **then**
- 11:     Sample  $a_c$  from  $\mathcal{C}$  using the distribution:
$$p(a_c) = \begin{cases} 0.5 & \text{if } a_c \in \mathcal{B}, \\ \frac{1}{2^{|\mathcal{B}|} |\mathcal{C} \setminus \mathcal{B}|} & \text{otherwise.} \end{cases}$$
- 12:     **end if**
- 13:  $d \leftarrow \arg \max_j u_{jc}$ , with ties broken randomly. Moreover, if there is a tie,  $d$  is not allowed to be equal to  $c$ .
- 14: Compare arms  $a_c$  and  $a_d$  and increment  $w_{cd}$  or  $w_{dc}$  depending on which arm wins.
- 15: **end for**

**Return:** An arm  $a_c$  that beats the most arms, i.e.,  $c$  with the largest count  $\# \left\{ j \mid \frac{w_{cj}}{w_{cj} + w_{jc}} > \frac{1}{2} \right\}$ .

---

# RUCB Algorithm: MergeRUCB

- Idea: partition the  $K$  arms into small batches and compares arms within each batch.

# RUCB Algorithm: MergeRUCB

- Idea: partition the  $K$  arms into small batches and compares arms within each batch.
- Regret:  $O(K \log T)$   
MergeRUCB does not require global pairwise comparisons between all pairs of arms than RUCB.

# Relative UCB (RUCB)

Reference: Reference: Relative Upper Confidence Bound for the K-Armed Dueling Bandit Problem, 2014



# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner**
- 5 Dependent Arms
- 6 Conclusion

Mainly generalized from traditional MAB algorithms along two lines:

- UCB-type: RUCB, CCB
- MED(minimal empirical divergence)-type: RMED

# SAVAGE Algorithm

- Motivation: solve the  $N$  dimensional parametric decision problem with high confidence.

- Motivation: solve the  $N$  dimensional parametric decision problem with high confidence.
- Idea:
  - work by reducing progressively a box-shaped confidence set  $H$  until a single decision remains in  $f(H)$ .
  - (more explicitly) eliminate the arms loss with high probability to another arm until leave with Condorcet(Copeland) winner.

- Motivation: solve the  $N$  dimensional parametric decision problem with high confidence.
- Idea:
  - work by reducing progressively a box-shaped confidence set  $H$  until a single decision remains in  $f(H)$ .
  - (more explicitly) eliminate the arms loss with high probability to another arm until leave with Condorcet(Copeland) winner.
- Regret:  $O(K^2 \log T)$  outperform IF and BTM for moderate  $K$

# SAVAGE Algorithm

---

**Algorithm 1** SAVAGE algorithm

---

```
1: Input:  $\mathbf{X} = (X_1, \dots, X_N)$ ,  $f$ ,  $\mathcal{F}$ ,  $T$ ,  $\delta$ 
2: Initialization:
3:  $\mathcal{W} := \{1, \dots, N\}$ ,  $\mathcal{H} := \mathcal{F}$ ,  $s := 1$ 
4:  $\forall i \in \mathcal{W} : \hat{\mu}_i := 1/2$ , and  $t_i := 0$ 
5: while  $\neg \text{Accept}(f, \mathcal{H}, \mathcal{W}) \wedge s \leq T$  do
6:   Pick a variable index  $i \in \arg \min_{\mathcal{W}} \{t_1, \dots, t_N\}$ 
7:    $t_i := t_i + 1$ 
8:   Sample the  $i^{\text{th}}$  distribution  $x_i \leftarrow X_i$ 
9:    $\hat{\mu}_i := (1 - \frac{1}{t_i})\hat{\mu}_i + \frac{1}{t_i}x_i$ 
10:   $\mathcal{H} := \mathcal{H} \cap \{\mathbf{x} \mid |x_i - \hat{\mu}_i| < c(t_i)\}$ 
11:   $\mathcal{W} := \mathcal{W} \setminus \{j \mid \text{IndepTest}(f, \mathcal{H}, j)\}$ 
12:   $s := s + 1$ 
13: end while
14: return  $\hat{d} \in f(\mathcal{H})$ 
```

---

# SAVAGE Algorithm

Reference: Reference: Generic Exploration and K-armed Voting Bandits, 2013

# Copeland Confidence Bounds (CCB)

- Motivation: design for small number of arms Copeland winner problem



# Copeland Confidence Bounds (CCB)

- Motivation: design for small number of arms Copeland winner problem
- Principle: optimism followed by pessimism
  - optimistic copeland winner: have confidence to beat other arms
  - pessimistic opponent: have confidence to beat potential winner

# Copeland Confidence Bounds (CCB)

- Motivation: design for small number of arms Copeland winner problem
- Principle: optimism followed by pessimism
  - optimistic copeland winner: have confidence to beat other arms
  - pessimistic opponent: have confidence to beat potential winner
- Regret:  $O(K^2 + K \log T)$

# Copeland Confidence Bounds (CCB)

---

**Algorithm 1** Copeland Confidence Bound

---

**Input:** A Copeland dueling bandit problem and an exploration parameter  $\alpha > \frac{1}{2}$ .

- 1:  $\mathbf{W} = [w_{ij}] \leftarrow \mathbf{0}_{K \times K}$  // 2D array of wins:  $w_{ij}$  is the number of times  $a_i$  beat  $a_j$
  - 2:  $\mathcal{B}_1 = \{a_1, \dots, a_K\}$  // potential best arms
  - 3:  $\mathcal{B}_1^i = \emptyset$  for each  $i = 1, \dots, K$  // potential to beat  $a_i$
  - 4:  $\overline{L}_C = K$  // estimated max losses of a Copeland winner
  - 5: **for**  $t = 1, 2, \dots$  **do**
  - 6:    $\mathbf{U} := [u_{ij}] = \frac{\mathbf{W}}{\mathbf{W} + \mathbf{W}^T} + \sqrt{\frac{\alpha \ln t}{\mathbf{W} + \mathbf{W}^T}}$  and  $\mathbf{L} := [l_{ij}] = \frac{\mathbf{W}}{\mathbf{W} + \mathbf{W}^T} - \sqrt{\frac{\alpha \ln t}{\mathbf{W} + \mathbf{W}^T}}$ , with  $u_{ii} = l_{ii} = \frac{1}{2}, \forall i$
  - 7:    $\overline{\text{Cpld}}(a_i) = \# \{k \mid u_{ik} \geq \frac{1}{2}, k \neq i\}$  and  $\text{Cpld}(a_i) = \# \{k \mid l_{ik} \geq \frac{1}{2}, k \neq i\}$
  - 8:    $\mathcal{C}_t = \{a_i \mid \overline{\text{Cpld}}(a_i) = \max_j \overline{\text{Cpld}}(a_j)\}$
  - 9:   Set  $\mathcal{B}_t \leftarrow \mathcal{B}_{t-1}$  and  $\mathcal{B}_t^i \leftarrow \mathcal{B}_{t-1}^i$  and update as follows:
    - A. Reset disproven hypotheses:** If for any  $i$  and  $a_j \in \mathcal{B}_t^i$  we have  $l_{ij} > 0.5$ , reset  $\mathcal{B}_t$ ,  $\overline{L}_C$  and  $\mathcal{B}_t^k$  for all  $k$  (i.e. set them to their original values as in Lines 2–4 above).
    - B. Remove non-Copeland winners:** For each  $a_i \in \mathcal{B}_t$ , if  $\overline{\text{Cpld}}(a_i) < \text{Cpld}(a_j)$  holds for any  $j$ , set  $\mathcal{B}_t \leftarrow \mathcal{B}_t \setminus \{a_i\}$ , and if  $|\mathcal{B}_t^i| \neq \overline{L}_C + 1$ , then set  $\mathcal{B}_t^i \leftarrow \{a_k \mid u_{ik} < 0.5\}$ . However, if  $\mathcal{B}_t = \emptyset$ , reset  $\mathcal{B}_t$ ,  $\overline{L}_C$  and  $\mathcal{B}_t^k$  for all  $k$ .
    - C. Add Copeland winners:** For any  $a_i \in \mathcal{C}_t$  with  $\overline{\text{Cpld}}(a_i) = \text{Cpld}(a_i)$ , set  $\mathcal{B}_t \leftarrow \mathcal{B}_t \cup \{a_i\}$ ,  $\mathcal{B}_t^i \leftarrow \emptyset$  and  $\overline{L}_C \leftarrow K - 1 - \overline{\text{Cpld}}(a_i)$ . For each  $j \neq i$ , if we have  $|\mathcal{B}_t^j| < \overline{L}_C + 1$ , set  $\mathcal{B}_t^j \leftarrow \emptyset$ , and if  $|\mathcal{B}_t^j| > \overline{L}_C + 1$ , randomly choose  $\overline{L}_C + 1$  elements of  $\mathcal{B}_t^j$  and remove the rest.
  - 10: With probability  $1/4$ , sample  $(c, d)$  uniformly from the set  $\{(i, j) \mid a_j \in \mathcal{B}_t^i \text{ and } 0.5 \in [l_{ij}, u_{ij}]\}$  (if it is non-empty) and skip to Line 14.
  - 11: If  $\mathcal{B}_t \cap \mathcal{C}_t \neq \emptyset$ , then with probability  $2/3$ , set  $\mathcal{C}_t \leftarrow \mathcal{B}_t \cap \mathcal{C}_t$ .
  - 12: Sample  $a_c$  from  $\mathcal{C}_t$  uniformly at random.
  - 13: With probability  $1/2$ , choose the set  $\mathcal{B}^i$  to be either  $\mathcal{B}_t^i$  or  $\{a_1, \dots, a_K\}$  and then set  $d \leftarrow \arg \max_{\{j \in \mathcal{B}^i \mid l_{jc} \leq 0.5\}} u_{jc}$ . If there is a tie,  $d$  is not allowed to be equal to  $c$ .
  - 14: Compare arms  $a_c$  and  $a_d$  and increment  $w_{cd}$  or  $w_{dc}$  depending on which arm wins.
  - 15: **end for**
-

# Scalable Confidence bounds (SCB)

- Motivation: design for large-scale arms Coepland winner problem

# Scalable Confidence bounds (SCB)

- Motivation: design for large-scale arms Copeland winner problem
- Algorithm difference from CCB:
  - Approximate Copeland Solver
  - KL-based arm-elimination identification

# Scalable Confidence bounds (SCB)

- Motivation: design for large-scale arms Copeland winner problem
- Algorithm difference from CCB:
  - Approximate Copeland Solver
  - KL-based arm-elimination identification
- Regret:  $O(K \log K \log T)$

# Scalable Confidence Bounds (SCB)

---

## Algorithm 3 Scalable Copeland Bandits

---

**Input:** A Copeland dueling bandit problem with preference matrix  $\mathbf{P} = [p_{ij}]$

---

- 1: **for all**  $r = 1, 2, \dots$  **do**
  - 2:   Set  $T = 2^{2^r}$  and run Algorithm 2 with failure probability  $\log(T)/T$  in order to find an exact Copeland winner ( $\epsilon = 0$ ); force-terminate if it requires more than  $T$  queries.
  - 3:   Let  $T_0$  be the number of queries used by invoking Algorithm 2 and let  $a_i$  be the arm produced by it; query the pair  $(a_i, a_i)$   $T - T_0$  times.
  - 4: **end for**
-

Reference: Copeland Dueling Bandits, 2015



# Relative Minimal Empirical Divergence (RMED)

- Idea: Minimal Empirical Divergence

# Relative Minimal Empirical Divergence (RMED)

- Idea: Minimal Empirical Divergence
- Theory: Empirical Divergence  $I(t) = \sum_{\{j | p_{ij} < .5\}} d(\hat{p}_{ij}, .5) N_{ij}(t)$ :
  - defines the minimum number of comparisons to prove  $i$  is inferior to  $j$
  - likelihood function of Condorcet winner:  $\exp(-I(t))$

# Relative Minimal Empirical Divergence (RMED)

- Idea: Minimal Empirical Divergence
- Theory: Empirical Divergence  $I(t) = \sum_{\{j|p_{ij}<.5\}} d(\hat{p}_{ij}, .5)N_{ij}(t)$ :
  - defines the minimum number of comparisons to prove  $i$  is inferior to  $j$
  - likelihood function of Condorcet winner:  $\exp(-I(t))$
- Algorithm: RMED1, RMED2, RMED2FH  
(Unlike RMED1, RMED2 keeps drawing pairs of arms  $i, j$  at least  $\alpha \log \log t$  times (Line 10 in Algorithm 1))

# Relative Minimal Empirical Divergence (RMED)

- Idea: Minimal Empirical Divergence
- Theory: Empirical Divergence  $l(t) = \sum_{\{j|p_{ij}<.5\}} d(\hat{p}_{ij}, .5)N_{ij}(t)$ :
  - defines the minimum number of comparisons to prove  $i$  is inferior to  $j$
  - likelihood function of Condorcet winner:  $\exp(-l(t))$
- Algorithm: RMED1, RMED2, RMED2FH  
(Unlike RMED1, RMED2 keeps drawing pairs of arms  $i, j$  at least  $\alpha \log \log t$  times (Line 10 in Algorithm 1))
- Regret:  $R_T \geq \sum_{k=2}^K \min_{\{j|p_{ij}<.5\}} \frac{(\Delta_{1i} + \Delta_{1j}) \log T}{2d(p_{ij}, .5)}$

# Relative Minimal Empirical Divergence (RMED)

---

**Algorithm 1** Relative Minimum Empirical Divergence (RMED) Algorithm

---

```
1: Input:  $K$  arms,  $f(K) \geq 0$ .  $\alpha > 0$  (RMED2FH, RMED2).  $T$  (RMED2FH).
2:  $L \leftarrow \begin{cases} 1 & \text{(RMED1, RMED2)} \\ \lceil \alpha \log \log T \rceil & \text{(RMED2FH)} \end{cases}$ .
3: Initial phase: draw each pair of arms  $L$  times. At the end of this phase,  $t = L(K-1)K/2$ .
4: if RMED2FH then
5:   For each arm  $i \in [K]$ , fix  $\hat{b}^*(i)$  by (6).
6: end if
7:  $L_C, L_R \leftarrow [K], L_N \leftarrow \emptyset$ .
8: while  $t \leq T$  do
9:   if RMED2 then
10:    Draw all pairs  $(i, j)$  until it reaches  $N_{i,j}(t) \geq \alpha \log \log t$ .  $t \leftarrow t + 1$  for each draw.
11:   end if
12:   for  $l(t) \in L_C$  in an arbitrarily fixed order do
13:    Select  $m(t)$  by using  $\begin{cases} \text{Algorithm 2} & \text{(RMED1)} \\ \text{Algorithm 3} & \text{(RMED2, RMED2FH)} \end{cases}$ .
14:    Draw arm pair  $(l(t), m(t))$ .
15:     $L_R \leftarrow L_R \setminus \{l(t)\}$ .
16:     $L_N \leftarrow L_N \cup \{j\}$  (without a duplicate) for any  $j \notin L_R$  such that  $\mathcal{J}_j(t)$  holds.
17:     $t \leftarrow t + 1$ .
18:   end for
19:    $L_C, L_R \leftarrow L_N, L_N \leftarrow \emptyset$ .
20: end while
```

---

# Relative Minimal Empirical Divergence (RMED)

Reference: Regret Lower Bound and Optimal Algorithm, 2015

# Double Thompson Sampling

- Motivation:
  - Extends the Thompson Sampling in relative setting to apply dueling bandit
  - TS achieves lower regret in practice and more robust than other algorithms

# Double Thompson Sampling

- Motivation:
  - Extends the Thompson Sampling in relative setting to apply dueling bandit
  - TS achieves lower regret in practice and more robust than other algorithms
- Idea: Sample from the parameters representing the number of winning:  $B_{ij}$



# Double Thompson Sampling

- Motivation:
  - Extends the Thompson Sampling in relative setting to apply dueling bandit
  - TS achieves lower regret in practice and more robust than other algorithms
- Idea: Sample from the parameters representing the number of winning:  $B_{ij}$
- Elimination:
  - RUCB-based: choose the first arm based on upper confidence bound
  - RLCB-based: choose the second arm only from the uncertain pair (with high lower confidence bound)
  - Objective: Avoid falling into the suboptimal comparisons.

# Double Thompson Sampling

- Motivation:
  - Extends the Thompson Sampling in relative setting to apply dueling bandit
  - TS achieves lower regret in practice and more robust than other algorithms
- Idea: Sample from the parameters representing the number of winning:  $B_{ij}$
- Elimination:
  - RUCB-based: choose the first arm based on upper confidence bound
  - RLCB-based: choose the second arm only from the uncertain pair (with high lower confidence bound)
  - Objective: Avoid falling into the suboptimal comparisons.
- Regret:  $O(K \log T + K^2 \log \log T)$

# Double Thompson Sampling

DT-S+

- change the selection of the first arms
- minimize the regret of comparison  $\tilde{R}_i^{(1)}(t)$

# Double Thompson Sampling

---

**Algorithm 1** D-TS for Copeland Dueling Bandits

---

```
1: Init:  $B \leftarrow \mathbf{0}_{K \times K}$ ; //  $B_{ij}$  is the number of time-slots that the user prefers arm  $i$  to  $j$ .
2: for  $t = 1$  to  $T$  do
3:   // Phase 1: Choose the first candidate  $a^{(1)}$ 
4:    $U := [u_{ij}]$ ,  $L := [l_{ij}]$ , where  $u_{ij} = \frac{B_{ij}}{B_{ij}+B_{ji}} + \sqrt{\frac{\alpha \log t}{B_{ij}+B_{ji}}}$ ,  $l_{ij} = \frac{B_{ij}}{B_{ij}+B_{ji}} - \sqrt{\frac{\alpha \log t}{B_{ij}+B_{ji}}}$ , if
      $i \neq j$ , and  $u_{ii} = l_{ii} = 1/2$ ,  $\forall i$ ; //  $\frac{x}{0} := 1$  for any  $x$ .
5:    $\hat{\zeta}_i \leftarrow \frac{1}{K-1} \sum_{j \neq i} \mathbb{1}(u_{ij} > 1/2)$ ; // Upper bound of the normalized Copeland score.
6:    $\mathcal{C} \leftarrow \{i : \hat{\zeta}_i = \max_j \hat{\zeta}_j\}$ ;
7:   for  $i, j = 1, \dots, K$  with  $i < j$  do
8:     Sample  $\theta_{ij}^{(1)} \sim \text{Beta}(B_{ij} + 1, B_{ji} + 1)$ ;
9:      $\theta_{ji}^{(1)} \leftarrow 1 - \theta_{ij}^{(1)}$ ;
10:  end for
11:   $a^{(1)} \leftarrow \arg \max_{i \in \mathcal{C}} \sum_{j \neq i} \mathbb{1}(\theta_{ij}^{(1)} > 1/2)$ ; // Choosing from  $\mathcal{C}$  to eliminate likely non-winner
     arms; Ties are broken randomly.
12:  // Phase 2: Choose the second candidate  $a^{(2)}$ 
13:  Sample  $\theta_{ia^{(1)}}^{(2)} \sim \text{Beta}(B_{ia^{(1)}} + 1, B_{a^{(1)}i} + 1)$  for all  $i \neq a^{(1)}$ , and let  $\theta_{a^{(1)}a^{(1)}}^{(2)} = 1/2$ ;
14:   $a^{(2)} \leftarrow \arg \max_{i: l_{ia^{(1)}} \leq 1/2} \theta_{ia^{(1)}}^{(2)}$ ; // Choosing only from uncertain pairs.
15:  // Compare and Update
16:  Compare pair  $(a^{(1)}, a^{(2)})$  and observe the result  $w$ ;
17:  Update  $B$ :  $B_{a^{(1)}a^{(2)}} \leftarrow B_{a^{(1)}a^{(2)}} + 1$  if  $w = 1$ , or  $B_{a^{(2)}a^{(1)}} \leftarrow B_{a^{(2)}a^{(1)}} + 1$  if  $w = 2$ ;
18: end for
```

---

# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner
- 5 Dependent Arms**
- 6 Conclusion

# Dependent Arms

- Motivation: solve the problem of dueling bandit arm with dependent arms

# Dependent Arms

- Motivation: solve the problem of dueling bandit arm with dependent arms
- Gaussian Process:
  - use covariance between arms to model the dependency
  - posterior inference updates the mean reward estimates for all arms and update the dependency

- Idea:
  - reduces the multi-dueling bandits problem to a conventional MAB bandit
  - simultaneously dueling multiple arms as well as modeling dependencies between arms using a kernel



- Idea:
  - reduces the multi-dueling bandits problem to a conventional MAB bandit
  - simultaneously dueling multiple arms as well as modeling dependencies between arms using a kernel
- Assumption: approximate linearity (using linear function to estimate the preference based on utility  $u$ )

- Idea:
  - reduces the multi-dueling bandits problem to a conventional MAB bandit
  - simultaneously dueling multiple arms as well as modeling dependencies between arms using a kernel
- Assumption: approximate linearity (using linear function to estimate the preference based on utility  $u$ )
- Algorithm: Ind-Self-Sparring, Kernel-Self-Sparring

---

## Algorithm 2 SELFSPARRING

---

**input** arms  $1, \dots, K$  in space  $S$ ,  $m$  the number of arms drawn at each iteration,  $\eta$  the learning rate

- 1: Set prior  $D_0$  over  $S$
  - 2: **for**  $t = 1, 2, \dots$  **do**
  - 3:   **for**  $j = 1, \dots, m$  **do**
  - 4:     select arm  $i_j(t)$  using  $D_{t-1}$
  - 5:   **end for**
  - 6:   Play  $m$  arms  $\{i_j(t)\}_j$  and observe  $m \times m$  pairwise feedback matrix  $R = \{r_{ij} \in \{0, 1, \emptyset\}\}_{m \times m}$
  - 7:   update  $D_{t-1}$  using  $R$  to obtain  $D_t$
  - 8: **end for**
-

## Ind-Self-Sparring:

- independent cases
- SBM: Beta-Bernoulli Thompson Sampling
- Regret:  $O(K \log T / \Delta)$

---

**Algorithm 3** INDSELFSPARRING

---

**input**  $m$  the number of arms drawn at each iteration,  $\eta$   
the learning rate

- 1: For each arm  $i = 1, 2, \dots, K$ , set  $S_i = 0, F_i = 0$ .
- 2: **for**  $t = 1, 2, \dots$  **do**
- 3:   **for**  $j = 1, \dots, m$  **do**
- 4:     For each arm  $i = 1, 2, \dots, K$ , sample  $\theta_i$  from  $Beta(S_i + 1, F_i + 1)$
- 5:     Select  $i_j(t) := \operatorname{argmax}_i \theta_i(t)$
- 6:   **end for**
- 7:   Play  $m$  arms  $\{i_j(t)\}_j$ , observe pairwise feedback matrix  $R = \{r_{jk} \in \{0, 1, \emptyset\}\}_{m \times m}$
- 8:   **for**  $j, k = 1, \dots, m$  **do**
- 9:     **if**  $r_{jk} \neq \emptyset$  **then**
- 10:        $S_j \leftarrow S_j + \eta \cdot r_{jk}, F_j \leftarrow F_j + \eta(1 - r_{jk})$
- 11:     **end if**
- 12:   **end for**
- 13: **end for**

---

- dependent cases

- dependent cases
- Idea:
  - use Gaussian processes  $GP((b), k(b, b'))$ . to model the preference function  $f(b)$
  - apply Bayesian update using  $(i_j(t), r_{jk})$  to obtain  $(\mu_t, \sigma_t)$

- dependent cases
- Idea:
  - use Gaussian processes  $GP((b), k(b, b'))$ . to model the preference function  $f(b)$
  - apply Bayesian update using  $(i_j(t), r_{jk})$  to obtain  $(\mu_t, \sigma_t)$
- Regret:  $O(K \log T)$  (A remaining opening problem)



---

**Algorithm 4** KERNELSELFSPARRING

---

**input** Input space  $S$ , GP prior  $(\mu_0, \sigma_0)$ ,  $m$  the number of arms drawn at each iteration

```
1: for  $t = 1, 2, \dots$  do
2:   for  $j = 1, \dots, m$  do
3:     Sample  $f_j$  from  $(\mu_{t-1}, \sigma_{t-1})$ 
4:     Select  $i_j(t) := \operatorname{argmax}_x f_j(x)$ 
5:   end for
6:   Play  $m$  arms  $\{i_j(t)\}_j$ , observe pairwise feedback matrix  $R = \{r_{jk} \in \{0, 1, \emptyset\}\}_{m \times m}$ 
7:   for  $j, k = 1, \dots, m$  do
8:     if  $r_{jk} \neq \emptyset$  then
9:       apply Bayesian update using  $(i_j(t), r_{jk})$  to obtain  $(\mu_t, \sigma_t)$ 
10:    end if
11:  end for
12: end for
```

---

Reference: Multi-dueling Bandits with Dependent Arms, 2017

## CoSpar Algorithm:

- the user can suggest improvements in the form of coactive feedback
- Reference: Reference Preference-Based Learning for Exoskeleton Gait Optimization, 2020

# Table of Contents

- 1 Motivation
- 2 Problem Setting
- 3 Condorcet Winner
- 4 Copeland Winner
- 5 Dependent Arms
- 6 Conclusion**


# Application Scenarios of Main Algorithms

- IF(Interleaved Filter, 2009): strong stochastic transitivity (**SST**)
- BTM (2011): relaxed stochastic transitivity (**RST**)
- UBDB (2014): reducing the Dueling Bandits problem to the **conventional Multi-Armed** Bandits problem
- RUCB (2014): Apply UCB in relative setting (by **upper bound**)
- CCB (2015): Using Upper bound to solve **Copeland Bandit** problem
- RMED (2015): **Minimum Empirical Divergence**
- DTS (2016): Extends the **Thompson Sampling** in relative setting to apply dueling bandit
- Self-Sparring (2017): Solve Dueling Bandit problem with **dependent arms**

# Performance of Main Algorithm

Algorithm	Regret
IF	$O(K \log T / \Delta_{\min})$
BTM	$O(\frac{\gamma^T K}{\Delta_{\min}} \log T)$
Sparring	$O(K \log T)$
RUCB	$O(K^2 + K \log T)$
SAVAGE	$O(K^2 \log T)$
CCB	$O(K^2 + K \log T)$
RMED	$R_T \geq \sum_{k=2}^K \min_{\{j   p_{ij} < .5\}} \frac{(\Delta_{1i} + \Delta_{1j}) \log T}{2d(p_{ij}, .5)}$
DTS	$O(K \log T + K^2 \log \log T)$
Self-Sparring	$O(K \log T / \Delta_{\min})$ ?

# Performance of Main Algorithm



- Upper Bound: RUCB, CCB
- Thompson Sampling: DTS
- Minimum Empirical Divergence: RMED
- Reduce to Conventional MAB: Sparring, SelfSparring

## Explore-Exploit Tradeoff



# Conclusion

- Elicits preference feedback
  - Motivated by human-centric personalization
  - Explore-exploit Tradeoff
- Algorithm Basis
  - Upper Bound
  - Thompson Sampling
  - Minimum Empirical Divergence
  - Reduce to Conventional MAB

## Ongoing research

- Personalized clinical treatment (High-Dimensional Multi-arms dueling bandit)
- Dependent Arms (regret bound)
- Apply more basic MAB algorithms in Dueling Bandit (e.g. Gradient Descent, Upper Bound in Dependent arms)
- Complex dueling mechanisms (Algorithms different with conventional MAB problem)

# Reference

- The K-armed Dueling Bandits Problem, Yisong Yue, Josef Broder, Robert Kleinberg and Thorsten Joachims, COLT 2009
- Beat the Mean Bandit, by Yisong Yue and Thorsten Joachims, ICML 2011
- Generic Exploration and K-armed Voting Bandits, Tanguy Urvoy, Fabrice Clerot, Raphael Feraud, Sami Naamane, 2013
- Relative Upper Confidence Bound for the K-armed Dueling Bandit Problem, Masrour Zoghi, Shimon Whiteson, Remi Munos and Maarten de Rijke, ICML, 2014
- Reducing Dueling Bandits to Cardinal Bandits, Nir Ailon, Zohar Karnin and Thorsten Joachims, ICML 2014
- Copeland Dueling Bandits, Masrour Zoghi, Zohar Karnin, Shimon Whiteson, and Maarten de Rijke, 2015
- Regret Lower Bound and Optimal Algorithm in Dueling Bandit Problem, Junpei Komiyama, Junya Honda, Hisashi Kashima, Hiroshi Nakagawa, 2015

- Double Thompson Sampling for Dueling Bandits, Huasen Wu, Xin Liu, NIPS 2016
- Multi-dueling Bandits with Dependent Arms, Yanan Sui, Vincent Zhuang, Joel W. Burdick, Yisong Yue, 2017
- Preference-Based Learning for Exoskeleton Gait Optimization, Maegan Tucker, Ellen Novoseller, Claudia Kann, Yanan Sui, Yisong Yue, Joel W. Burdick, and Aaron D. Ames, 2020